

Math642__HW5__FyonaSun

Fyona Sun

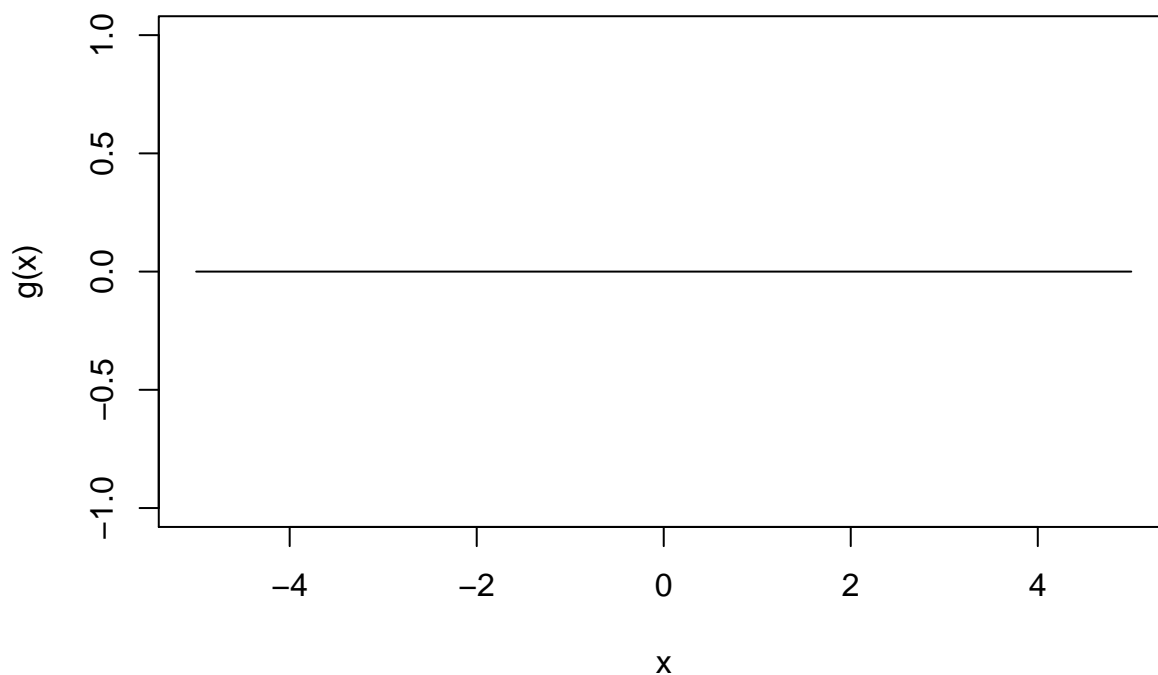
2/17/2020

7.2

a) $\lambda = \infty, m = 0$

To minimize \hat{g} when $\lambda = \infty$, $g(x)$ must equal to zero. $\hat{g}(x) = 0$

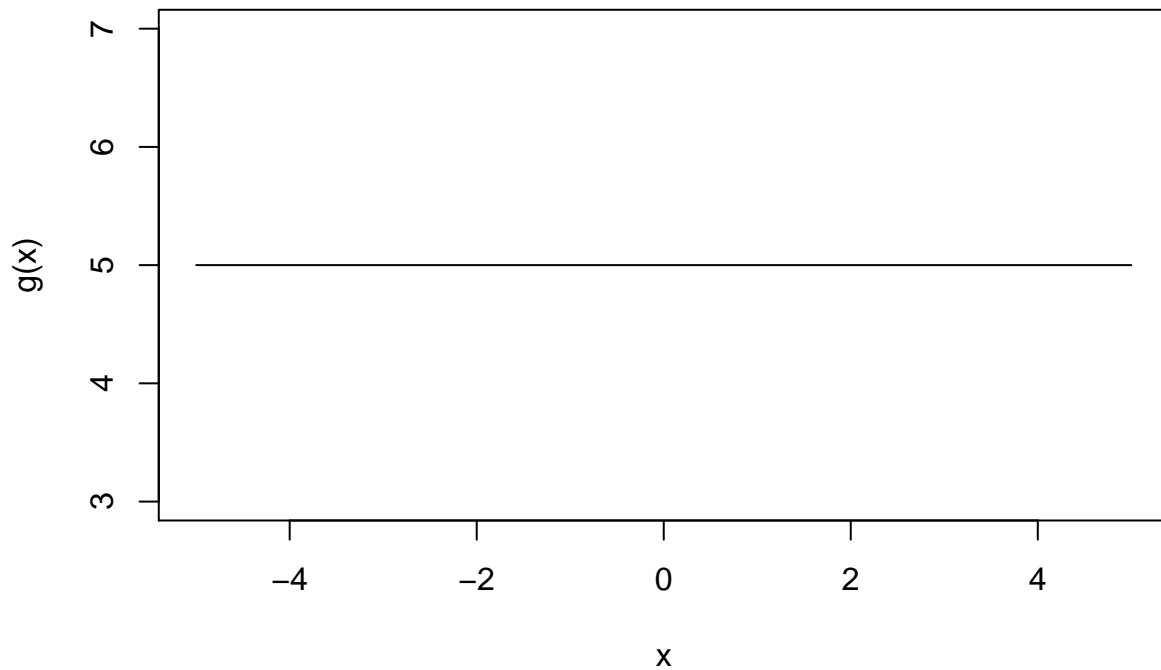
```
g<- function(x) {return(rep(0,length(x)))}  
x<- -5:5  
plot(x,g(x),type = 'l')
```



b) $\lambda = \infty, m = 1$

To minimize \hat{g} when $\lambda = \infty$, $g'(x)$ must equal to zero, that is $g(x)$ must be a constant. $\hat{g}(x) = c$, $c=5$ in the example.

```
g<- function(x) {return(rep(5,length(x)))}  
x<- -5:5  
plot(x,g(x),type = 'l')
```

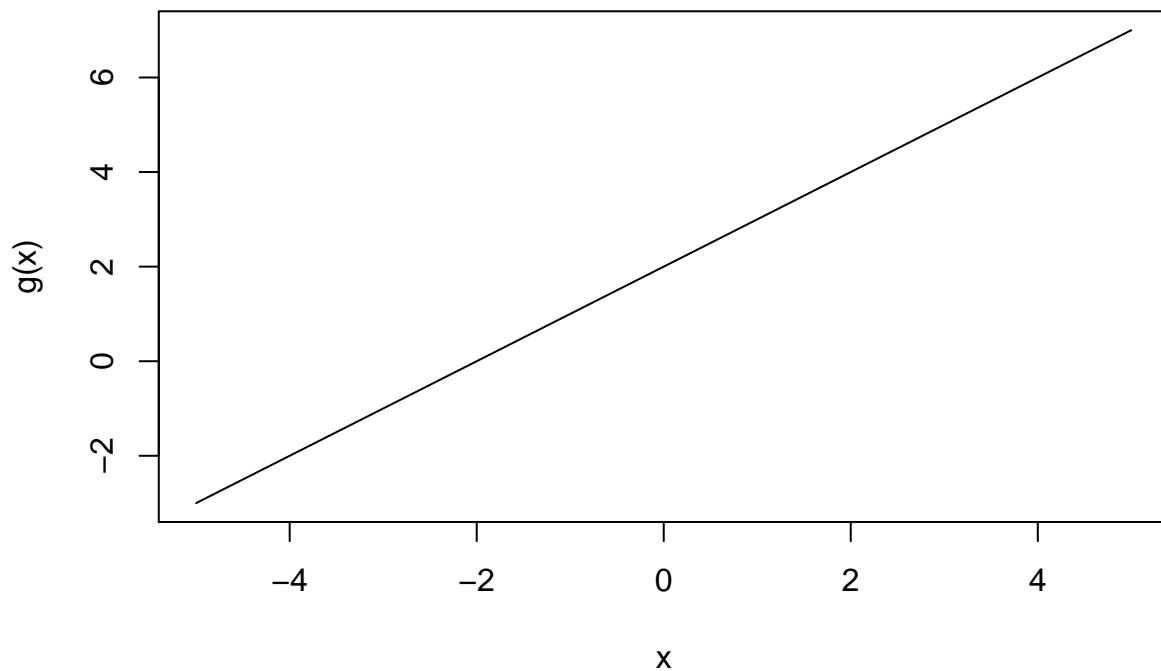


$\infty, m = 2$

c) $\lambda =$

To minimize \hat{g} when $\lambda = \infty$, $g''(x)$ must equal to zero, that is $g(x)$ must be a constant. $g(x) = ax + b$, $a=1$, $b=2$ in this example.

```
g<- function(x) {x+2}
x<- -5:5
plot(x,g(x),type = 'l')
```

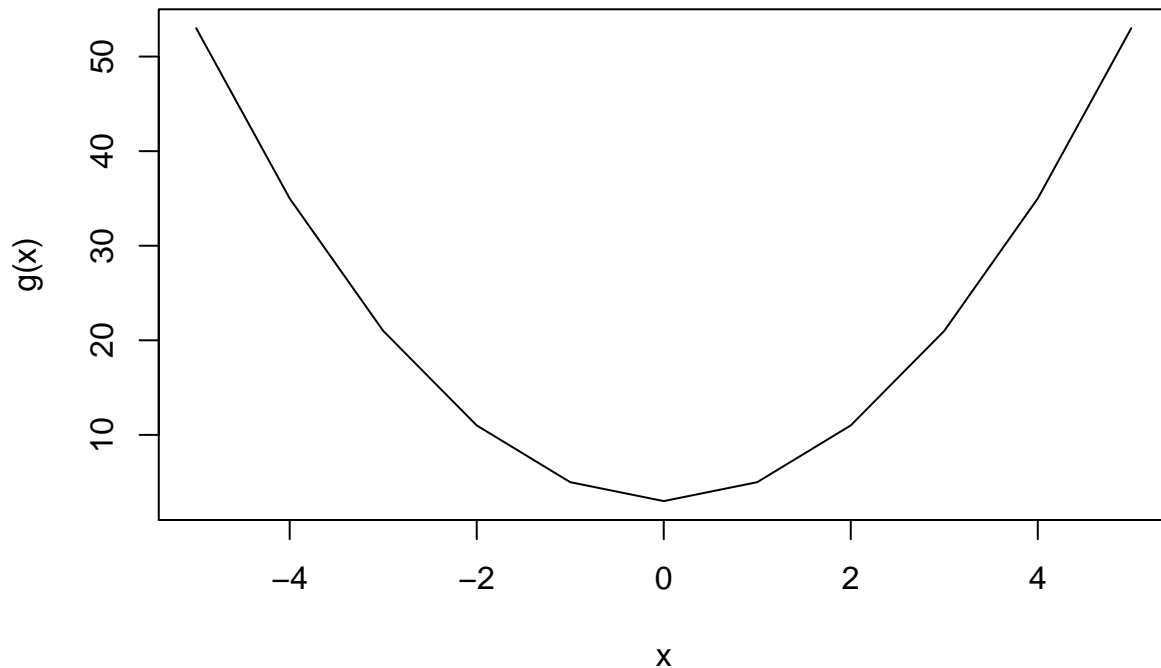


d) $\lambda =$

$\infty, m = 3$ To minimize \hat{g} when $\lambda = \infty$, $g'''(x)$ must equal to zero, that is $g(x)$ must be a constant. $g(x) = ax^2 + b$, $a=2$, $b=3$ in this example.

```
g<- function(x) {2*x^2+3}
x<- -5:5
```

```
plot(x,g(x),type = 'l')
```



0, $m = 3$ The penalty term doesn't have a function, so in this case g is the interpolating spline.

e) $\lambda =$

7.5

- Since \hat{g}_2 has a higher order of the penalty term, it has a higher order of polynomial. Compared with \hat{g}_1 , \hat{g}_2 is more flexible and probably has a smaller training RSS.
- Since \hat{g}_2 is more flexible, it could have over-fitting problem which leads to a higher testing RSS. Thus \hat{g}_1 would possibly have a smaller test RSS.
- When $\lambda = 0$, the penalty term disappears, that is $\hat{g}_1 = \hat{g}_2$. Therefore \hat{g}_1 and \hat{g}_2 would have the same training and test RSS.

7.9

This question uses the variables `dis` (the weighted mean of distances to five Boston employment centers) and `nox` (nitrogen oxides concentration in parts per 10 million) from the Boston data. We will treat `dis` as the predictor and `nox` as the response.

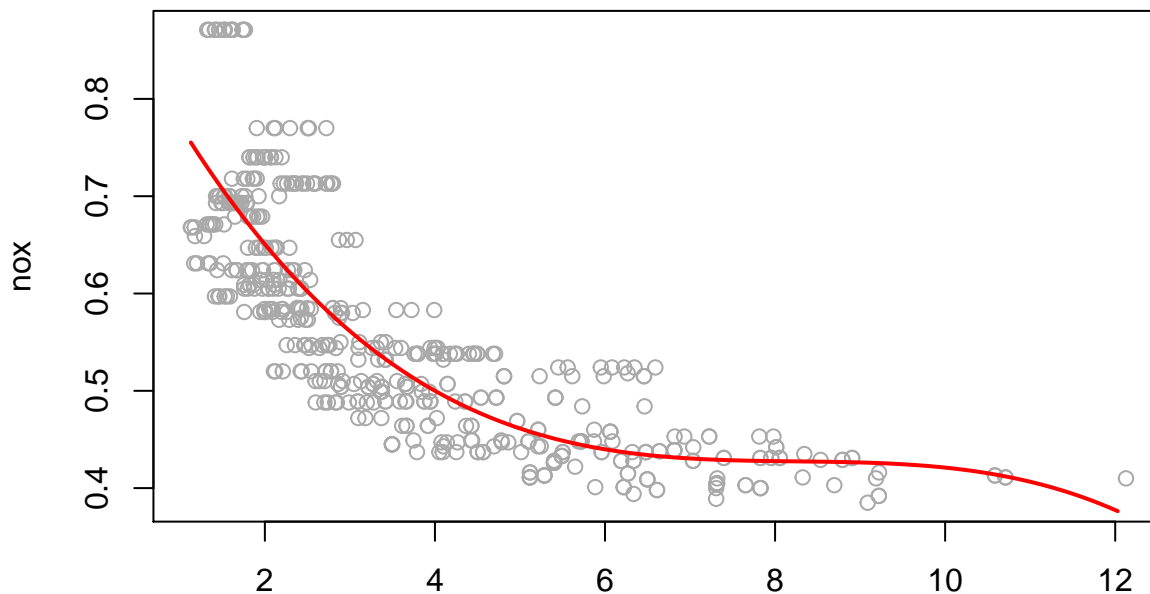
- Use the `poly()` function to fit a cubic polynomial regression to predict `nox` using `dis`. Report the regression output, and plot the resulting data and polynomial fits.

```
library(MASS)
attach(Boston)
set.seed(1)
fit <- lm(nox ~ poly(dis, 3), data = Boston)
summary(fit)
```

```
##
## Call:
## lm(formula = nox ~ poly(dis, 3), data = Boston)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.121130 -0.040619 -0.009738  0.023385  0.194904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.554695   0.002759  201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096   0.062071 -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330   0.062071  13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049   0.062071  -5.124 4.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06207 on 502 degrees of freedom
## Multiple R-squared:  0.7148, Adjusted R-squared:  0.7131
## F-statistic: 419.3 on 3 and 502 DF,  p-value: < 2.2e-16

dis.new <- seq(min(Boston$dis), max(Boston$dis), by = 0.1)
pred <- predict(fit, list(dis = dis.new))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.new, pred, col = "red", lwd = 2)
```



dis


From

the regression result, all the polynomial terms are significant. Thus $\hat{nox} = 0.554695 - 2.003096 * dis + 0.856330 * dis^2 - 0.318049 * dis^3$, which gives us the plotted line.

- (b) Plot the polynomial fits for a range of different polynomial degrees (say, from 1 to 10), and report the associated residual sum of squares.

```
rss <- rep(0, 10)
for (i in 1:10) {
  fit <- lm(nox ~ poly(dis, i), data = Boston)
  rss[i] <- sum(fit$residuals^2)
}
```

```
plot(1:10, rss, xlab = "Degree", ylab = "RSS", type = "l", col='red')
```




The RSS decreases as the degree of the polynomial decreases. It achieves its minimum at degree 10.

- (c) Perform cross-validation or another approach to select the optimal degree for the polynomial, and explain your results.

```
library(boot)
cv.error <- rep(0, 10)
for (i in 1:10) {
  fit <- glm(nox ~ poly(dis, i), data = Boston)
  cv.error[i] <- cv.glm(Boston, fit, K = 10)$delta[1]
}

plot(1:10, cv.error, xlab = "Degree", ylab = "Test Mean Squared Error", type = "l")
```



```
cv.error
```

```
## [1] 0.005536329 0.004077147 0.003899587 0.003862127 0.004298590
## [6] 0.005095283 0.013680327 0.005284520 0.013355413 0.004148996
```

At degree of 3, the test MSE is 0.003874862 and it's the smallest of the 10 degrees.

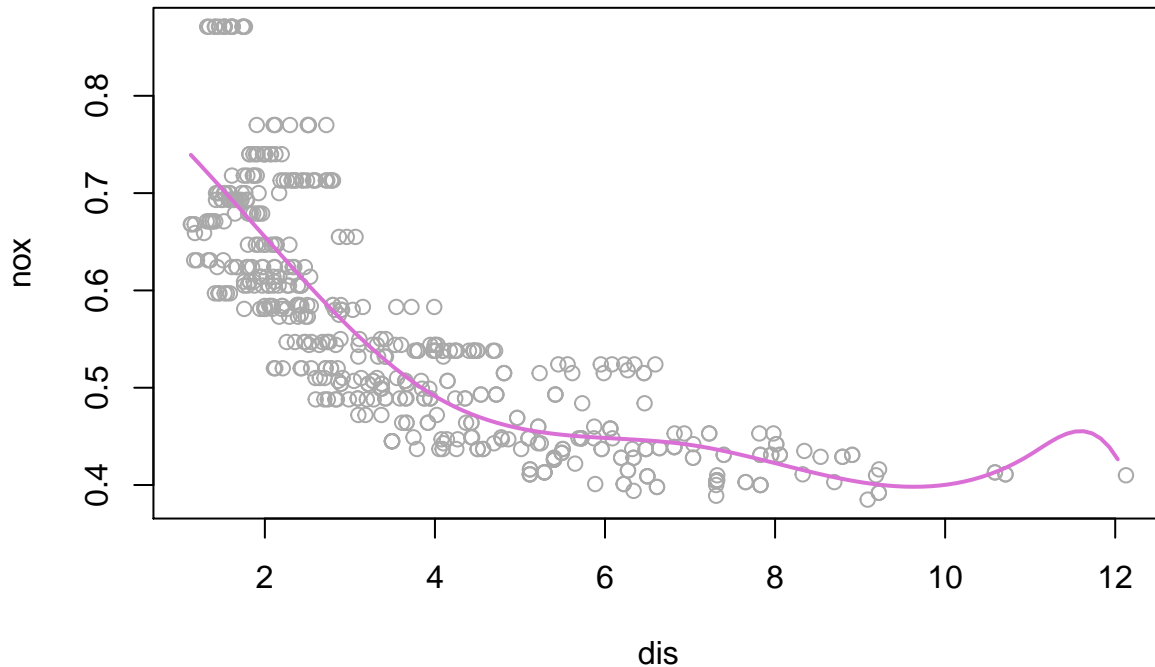
- (d) Use the `bs()` function to fit a regression spline to predict `nox` using `dis`. Report the output for the fit using four degrees of freedom. How did you choose the knots? Plot the resulting fit.

```
library(splines)
fit <- lm(nox ~ bs(dis, knots = c(4, 7, 11)), data = Boston)
summary(fit)

##
## Call:
## lm(formula = nox ~ bs(dis, knots = c(4, 7, 11)), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.124567 -0.040355 -0.008702  0.024740  0.192920
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      0.73926    0.01331  55.537  < 2e-16 ***
## bs(dis, knots = c(4, 7, 11))1 -0.08861    0.02504  -3.539  0.00044 ***
## bs(dis, knots = c(4, 7, 11))2 -0.31341    0.01680 -18.658  < 2e-16 ***
## bs(dis, knots = c(4, 7, 11))3 -0.26618    0.03147  -8.459  3.00e-16 ***
## bs(dis, knots = c(4, 7, 11))4 -0.39802    0.04647  -8.565  < 2e-16 ***
## bs(dis, knots = c(4, 7, 11))5 -0.25681    0.09001  -2.853  0.00451 **
## bs(dis, knots = c(4, 7, 11))6 -0.32926    0.06327  -5.204  2.85e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06185 on 499 degrees of freedom
```

```
## Multiple R-squared:  0.7185, Adjusted R-squared:  0.7151
## F-statistic: 212.3 on 6 and 499 DF,  p-value: < 2.2e-16


pred <- predict(fit, list(dis = dis.new))
plot(nox ~ dis, data = Boston, col = "darkgrey")
lines(dis.new, pred, col = "orchid", lwd = 2)
```



Here we have prespecified knots at `dis` 4, 7, and 11. This produces a spline with six basis functions. The regression indicates that all the spine terms are significant.

- (e) Now fit a regression spline for a range of degrees of freedom, and plot the resulting fits and report the resulting RSS. Describe the results obtained.

```
#Recall that a cubic spline with three knots has seven degrees of freedom
rss <- rep(0, 16)
for (i in 3:16) {
  fit <- lm(nox ~ bs(dis, df = i), data = Boston)
  rss[i] <- sum(fit$residuals^2)
}
plot(3:16, rss[-c(1, 2)], xlab = "Degrees of freedom", ylab = "RSS", type = "l")
```



```
rss[-c(1, 2)]
```

```
## [1] 1.934107 1.922775 1.840173 1.833966 1.829884 1.816995 1.825653
## [8] 1.792535 1.796992 1.788999 1.782350 1.781838 1.782798 1.783546
```

RSS decreases as the degrees of freedom increases at first, and RSS slightly increases after 14.

- (f) Perform cross-validation or another approach in order to select the best degrees of freedom for a regression spline on this data. Describe your results.

```
cv <- rep(0, 16)
for (i in 3:16) {
  fit <- glm(nox ~ bs(dis, df = i), data = Boston)
  cv[i] <- cv.glm(Boston, fit, K = 10)$delta[1]
}
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots =
## c(1.1296, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots =
## c(1.1296, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots =
## c(1.137, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots =
## c(1.137, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases
```

```
## Warning in bs(dis, degree = 3L, knots = c(`50%` = 3.0993), Boundary.knots =
## c(1.1296, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases
```



```

## Warning in bs(dis, degree = 3L, knots = c(`50%` = 3.0993), Boundary.knots =
## c(1.1296, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`50%` = 3.1523), Boundary.knots =
## c(1.1691, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`50%` = 3.1523), Boundary.knots =
## c(1.1691, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`33.3333%` = 2.38403333333333, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`33.3333%` = 2.38403333333333, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`33.3333%` = 2.38876666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`33.3333%` = 2.38876666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`25%` = 2.102875, `50%` =
## 3.26745, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`25%` = 2.102875, `50%` =
## 3.26745, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`25%` = 2.0788, `50%` = 3.2721, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`25%` = 2.0788, `50%` = 3.2721, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`20%` = 1.94264, `40%` =
## 2.62334, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`20%` = 1.94264, `40%` =
## 2.62334, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`20%` = 1.97036, `40%` =
## 2.66262, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`20%` = 1.97036, `40%` =
## 2.66262, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`16.6667%` = 1.87351666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

```

```

## Warning in bs(dis, degree = 3L, knots = c(`16.66667%` = 1.873516666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`16.66667%` = 1.844766666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`16.66667%` = 1.844766666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`14.28571%` = 1.79777142857143, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`14.28571%` = 1.79777142857143, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`14.28571%` = 1.79078571428571, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`14.28571%` = 1.79078571428571, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`12.5%` = 1.734325, `25%` =
## 2.0493, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`12.5%` = 1.734325, `25%` =
## 2.0493, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`12.5%` = 1.7275, `25%` =
## 2.0581, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`12.5%` = 1.7275, `25%` =
## 2.0581, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`11.11111%` = 1.653444444444444, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`11.11111%` = 1.653444444444444, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`11.11111%` = 1.718066666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`11.11111%` = 1.718066666666667, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`10%` = 1.6424, `20%` =
## 1.96376, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`10%` = 1.6424, `20%` =
## 1.96376, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`10%` = 1.64186, `20%` =

```

```

## 1.95434, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`10%` = 1.64186, `20%` =
## 1.95434, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`9.090909%` = 1.59590909090909, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`9.090909%` = 1.59590909090909, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`9.090909%` = 1.61450909090909, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`9.090909%` = 1.61450909090909, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`8.333333%` = 1.591425,
## `16.66667%` = 1.86565, : some 'x' values beyond boundary knots may cause
## ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`8.333333%` = 1.591425,
## `16.66667%` = 1.86565, : some 'x' values beyond boundary knots may cause
## ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`8.333333%` = 1.58948333333333, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`8.333333%` = 1.58948333333333, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`7.692308%` = 1.5888, `15.38462%`
## = 1.8172, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`7.692308%` = 1.5888, `15.38462%`
## = 1.8172, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c(`7.142857%` = 1.5311, `14.28571%`
## = 1.80062857142857, : some 'x' values beyond boundary knots may cause ill-
## conditioned bases

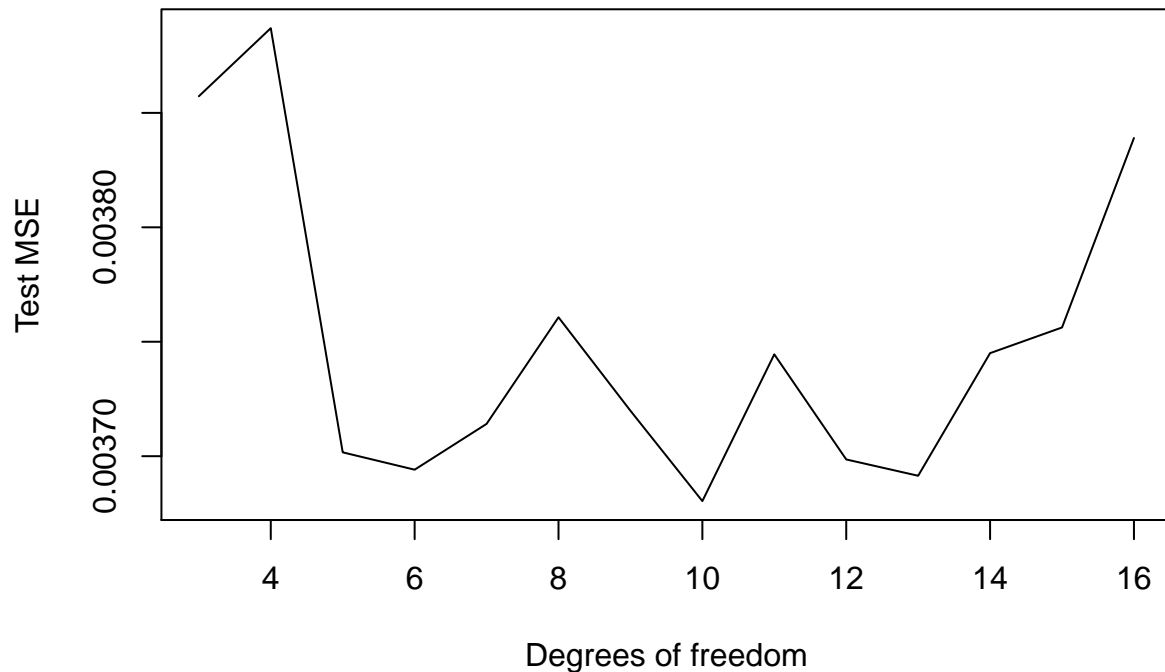
## Warning in bs(dis, degree = 3L, knots = c(`7.142857%` = 1.5311, `14.28571%`
## = 1.80062857142857, : some 'x' values beyond boundary knots may cause ill-
## conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`7.142857%` = 1.584, `14.28571%`
## = 1.81652857142857, : some 'x' values beyond boundary knots may cause ill-
## conditioned bases

## Warning in bs(dis, degree = 3L, knots = c(`7.142857%` = 1.584, `14.28571%`
## = 1.81652857142857, : some 'x' values beyond boundary knots may cause ill-
## conditioned bases

```

```
plot(3:16, cv[-c(1, 2)], xlab = "Degrees of freedom", ylab = "Test MSE", type = "l")
```



The test MSE minized at 8 degrees of freedom. ##7.11 (a) Generate a response Y and two predictors X1 and X2, with n = 100.

```
set.seed(1)
Y <- rnorm(100)
X1 <- rnorm(100)
X2 <- rnorm(100)
```

(b) Initialize β_1 to take on a value of your choice. It does not matter what value you choose.

```
beta1 <- 0.25
```

(c) Keeping β_1 fixed, fit the model $Y - \hat{\beta}_1 X_1 = \beta_0 + \beta_2 X_2 + \epsilon$

```
a<- Y-beta1*X1
beta2<- lm(a~X2)$coef[2]
```

```
beta2
```

```
##           X2
## 0.02743328
```

(d) Keeping β_2 fixed, fit the model $Y - \hat{\beta}_2 X_2 = \beta_0 + \beta_1 X_1 + \epsilon$

```
a<- Y-beta2*X2
beta1<- lm(a~X1)$coef[2]
```

```
beta1
```

```
##           X1
## 0.0005349403
```

(e) Write a for loop to repeat (c) and (d) 1,000 times.

```

iter<- 1000
df<- data.frame(0, 0.25, 0)
names(df)=c('beta0','beta1','beta2')
for (i in 1:iter) {
  beta1 <- df[nrow(df), 2]
  a <- Y - beta1 * X1
  beta2 <- lm(a ~ X2)$coef[2]
  a <- Y - beta2 * X2
  beta1 <- lm(a ~ X1)$coef[2]
  beta0 <- lm(a ~ X1)$coef[1]
  df[nrow(df) + 1,] <- list(beta0, beta1, beta2)
}

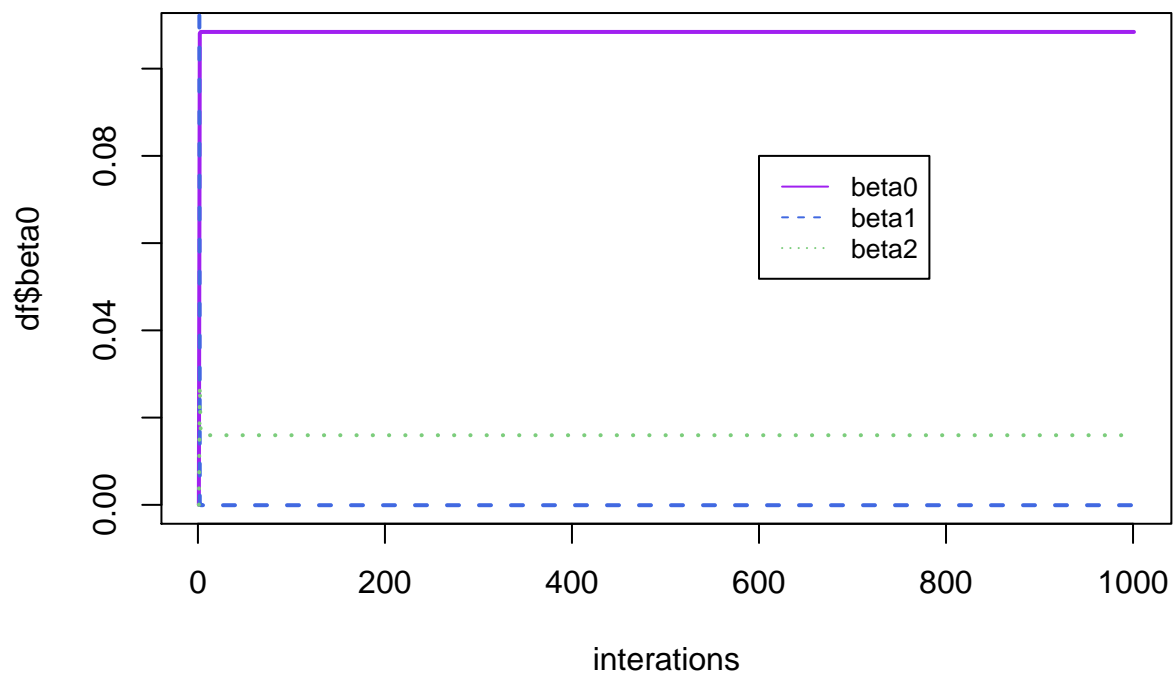
head(df)

##      beta0      beta1      beta2
## 1 0.0000000 2.500000e-01 0.0000000
## 2 0.1080935 5.349403e-04 0.02743328
## 3 0.1084100 -7.720619e-05 0.01598840
## 4 0.1084108 -7.870830e-05 0.01596032
## 5 0.1084108 -7.871198e-05 0.01596025
## 6 0.1084108 -7.871199e-05 0.01596025

plot(df$beta0, col = 'purple', type = 'l',lwd=2,xlab = 'iterations', main='Backfitting')
lines(df$beta1, col = 'royalblue',lwd=2, lty=2)
lines(df$beta2, col = 'palegreen3',lwd=2,lty=3)
legend(600,0.08,legend=c('beta0', 'beta1', 'beta2'),
      col=c('purple', 'royalblue','palegreen3'),lty=1:3,cex=0.8)

```

Backfitting

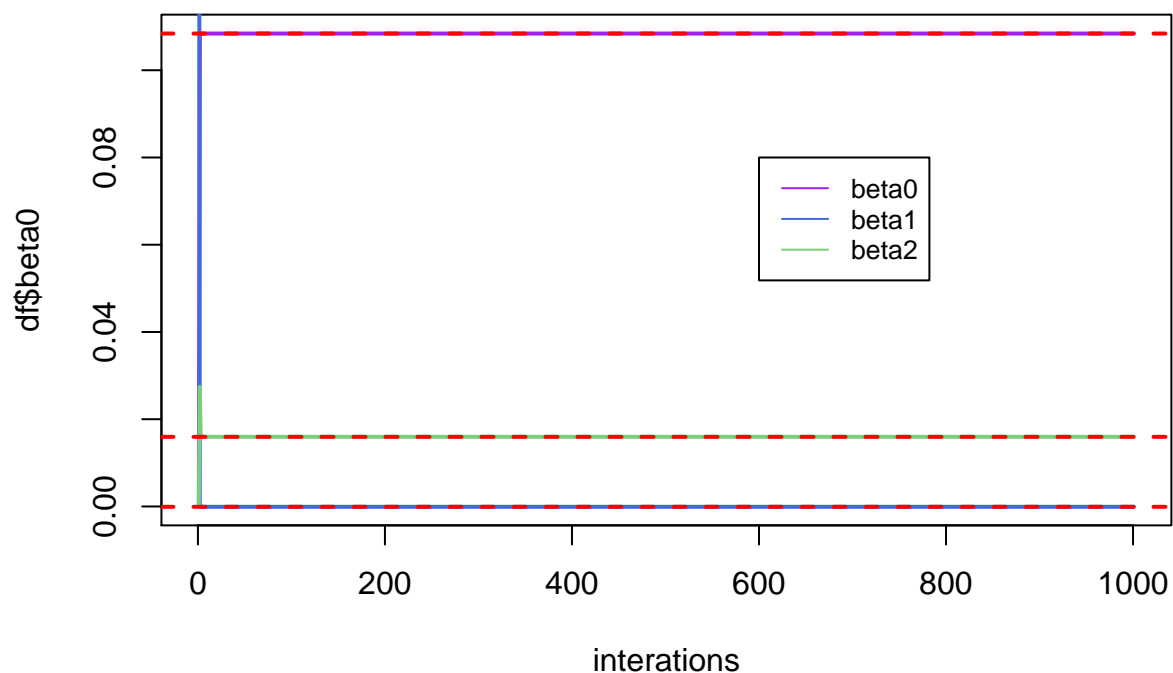


beta0, beta1 and beta2 attained the least squared values very quick and remained unchange.

- (f) Compare your answer in (e) to the results of simply performing multiple linear regression to predict Y using X1 and X2. Use the `abline()` function to overlay those multiple linear regression coefficient estimates on the plot obtained in (e).

```
lm.fit<-coef(lm(Y~X1+X2))
plot(df$beta0, col = 'purple', type='l',lwd=2,xlab = 'iterations', main='Backfitting vs Multiple Linear Regression')
lines(df$beta1, col = 'royalblue',lwd=2)
lines(df$beta2, col = 'palegreen3',lwd=2)
legend(600,0.08,legend=c('beta0', 'beta1', 'beta2'),
      col=c('purple', 'royalblue', 'palegreen3'),lty=1,cex=0.8)
abline(h=lm.fit[1], col="red", lty=2, lwd=2)
abline(h=lm.fit[2], col="red", lty=2, lwd=2)
abline(h=lm.fit[3], col="red", lty=2, lwd=2)
```

Backfitting vs Multiple Linear Regression



The red line on this graph show that the coefficients given by multiple regression match with the coefficients obtained by the backfitting estimation.

- (g) On this data set, how many backfitting iterations were required in order to obtain a “good” approximation to the multiple regression coefficient estimates?

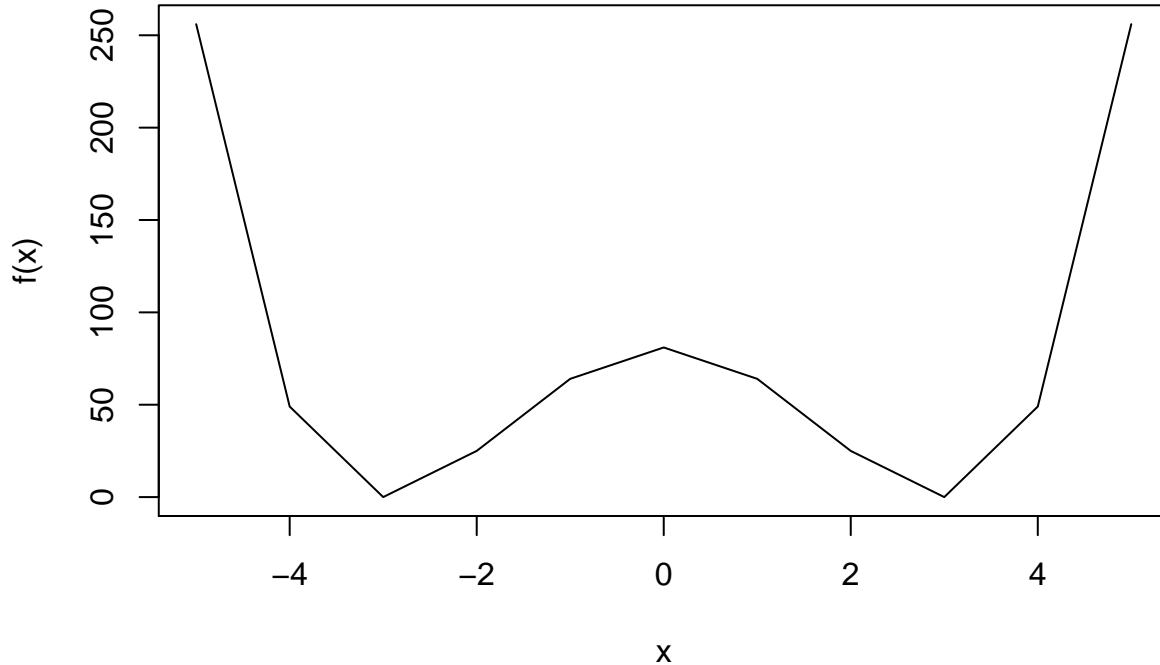
```
head(df)
```

```
##      beta0      beta1      beta2
## 1 0.000000 2.500000e-01 0.0000000
## 2 0.1080935 5.349403e-04 0.02743328
## 3 0.1084100 -7.720619e-05 0.01598840
## 4 0.1084108 -7.870830e-05 0.01596032
## 5 0.1084108 -7.871198e-05 0.01596025
## 6 0.1084108 -7.871199e-05 0.01596025
```

From the iteration results, beta0 converges to constant at iteration 4, beta1 converges at iteration4 and beta2 converges at iteration 5. So for this data set at least 5 iterations are required to obtain a ‘good’ approximation to the multiple regression coefficient estimates.

Write a Gradient Descent program in R to find the minimum of the equation. $y_i = (x_i^2 - 9)^2$

```
#plot the function
f<- function(x) ((x^2)-9)^2
x<- -5:5
plot(x,f(x),type = 'l')
```



```
#using gradient descent to find the numerical solution
```

```
gd<- function(x,f,lr,iter){
  x<- 0.001
  xtrace<- x
  ytrace<- f(x)
  lr<- 0.01
  for (i in 1:iter){
    delta<- (f(x+lr)-f(x))/lr
    x<- x - lr*delta
    xtrace<- c(xtrace,x)
    ytrace<- c(ytrace,f(x))
  }
  print(x)
  print(xtrace)
  print(ytrace)
}

iter<- 100
gd(x,f,lr,iter)
```

```
## [1] 2.994996
## [1] 0.001000000 0.003159985 0.006097550 0.010092603 0.015525787
## [6] 0.022914704 0.032963099 0.046627588 0.065207964 0.090468918
## [11] 0.124802828 0.171444235 0.234744257 0.320500755 0.436301219
## [16] 0.591731538 0.798054724 1.066438122 1.402955208 1.798175958
## [21] 2.212800507 2.574864061 2.816780641 2.933887813 2.976550286
```

```
## [26] 2.989708809 2.993505399 2.994577709 2.994878695 2.994963029
## [31] 2.994986648 2.994993261 2.994995113 2.994995632 2.994995777
## [36] 2.994995818 2.994995829 2.994995832 2.994995833 2.994995833
## [41] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [46] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [51] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [56] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [61] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [66] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [71] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [76] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [81] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [86] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [91] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [96] 2.994995833 2.994995833 2.994995833 2.994995833 2.994995833
## [101] 2.994995833
## [1] 8.099998e+01 8.099982e+01 8.099933e+01 8.099817e+01 8.099566e+01
## [6] 8.099055e+01 8.098044e+01 8.096087e+01 8.092348e+01 8.085274e+01
## [11] 8.071988e+01 8.047179e+01 8.001115e+01 7.916158e+01 7.760978e+01
## [16] 7.481997e+01 6.994159e+01 6.182220e+01 4.944504e+01 3.325325e+01
## [21] 1.683883e+01 5.617256e+00 1.135816e+00 1.539011e-01 1.964157e-02
## [26] 3.799642e-03 1.515189e-03 1.056532e-03 9.425885e-04 9.118258e-04
## [31] 9.033018e-04 9.009221e-04 9.002563e-04 9.000700e-04 9.000178e-04
## [36] 9.000032e-04 8.999991e-04 8.999979e-04 8.999976e-04 8.999975e-04
## [41] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [46] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [51] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [56] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [61] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [66] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [71] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [76] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [81] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [86] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [91] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [96] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [101] 8.999975e-04
```

The gradient descent with the starting $x=0.01$ and 100 iteration gives $x=2.994996$ at its minimum, which is close to one of the solution $x=3$.

#using gradient descent to find the numerical solution

```
gd<- function(x,f,lr,iter){
x<- -5
xtrace<- x
ytrace<- f(x)
lr<- 0.01
for (i in 1:iter){
  delta<- (f(x+lr)-f(x))/lr
  x<- x - lr*delta
  xtrace<- c(xtrace,x)
  ytrace<- c(ytrace,f(x))
}
print(x)
```



```

print(xtrace)
print(ytrace)
}

iter<- 100
gd(x,f,lr,iter)

```

```

## [1] -3.004996
## [1] -5.000000 -1.813180 -2.227648 -2.588589 -2.828867 -2.944730 -2.986823
## [8] -2.999789 -3.003528 -3.004584 -3.004880 -3.004964 -3.004987 -3.004993
## [15] -3.004995 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [22] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [29] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [36] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [43] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [50] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [57] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [64] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [71] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [78] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [85] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [92] -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996 -3.004996
## [99] -3.004996 -3.004996 -3.004996
## [1] 2.560000e+02 3.263127e+01 1.630208e+01 5.286345e+00 9.950248e-01
## [6] 1.079534e-01 6.223686e-03 1.607869e-06 4.486260e-04 7.576521e-04
## [11] 8.588826e-04 8.883865e-04 8.967386e-04 8.990844e-04 8.997418e-04
## [16] 8.999259e-04 8.999775e-04 8.999919e-04 8.999959e-04 8.999971e-04
## [21] 8.999974e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [26] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [31] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [36] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [41] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [46] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [51] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [56] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [61] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [66] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [71] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [76] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [81] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [86] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [91] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [96] 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04 8.999975e-04
## [101] 8.999975e-04

```

The gradient descent with the starting $x=-5$ and 100 iteration gives $x=-3.004996$ at its minimum, which is close to the other solution $x=-3$.